



## **Introduction to SAGA-BigJob**

### **Piloting Many Simulations on Many Supercomputers**

**RADICAL TEAM**

<http://radical.rutgers.edu>

# “Many Simulations” Scenarios: Context

- A Single “Application” is broken into many smaller tasks
  - Naturally decomposed or “by design” (algorithmic and infrastructural)
  - Varying task duration: hours to days to weeks
  - Often these tasks are “coupled”:
    - General concept for data sharing, synchronization, other dependency
- Coupling between these tasks
  - Uncoupled Tasks, Loosely-Coupled Tasks, Sequential dependencies..
  - Homogenous or Heterogeneous
  - Varying rate of coupling between tasks
  - Regular versus Irregular synchronization:
    - Temporal, Spatial
  - Ad hoc (pair-wise) exchange
    - No a priori determined exchange partners

# Common Characteristics and Requirements

- Many requirements, but focus on two common and critical requirements
- R1: Adaptive Application Formulation
  - Flexible composition
- R2: Dynamic Resource Utilization
  - Growing/shinking resource pool
- Need abstractions that:
  - A1: Decouple workload and resource management
  - A2: Provide these capabilities in an extensible and interoperable way
- Need to provide these abstractions as well engineered tools
  - Employ CI best practices
  - Software sustainability
- Is there a single simple tool that can support the above requirements across all machines on XSEDE, OSG independent of application type?

# Outline

## PART-I

- Scalable approaches to many simulations on many machines
  - Introducing the Pilot Abstraction
- SAGA-BigJob: A simple, extensible and interoperable PJS
- SAGA: The Interoperability Layer
- BigJob: Case Studies of varying the coupling between tasks

## PART-II

- Hands on Tutorial

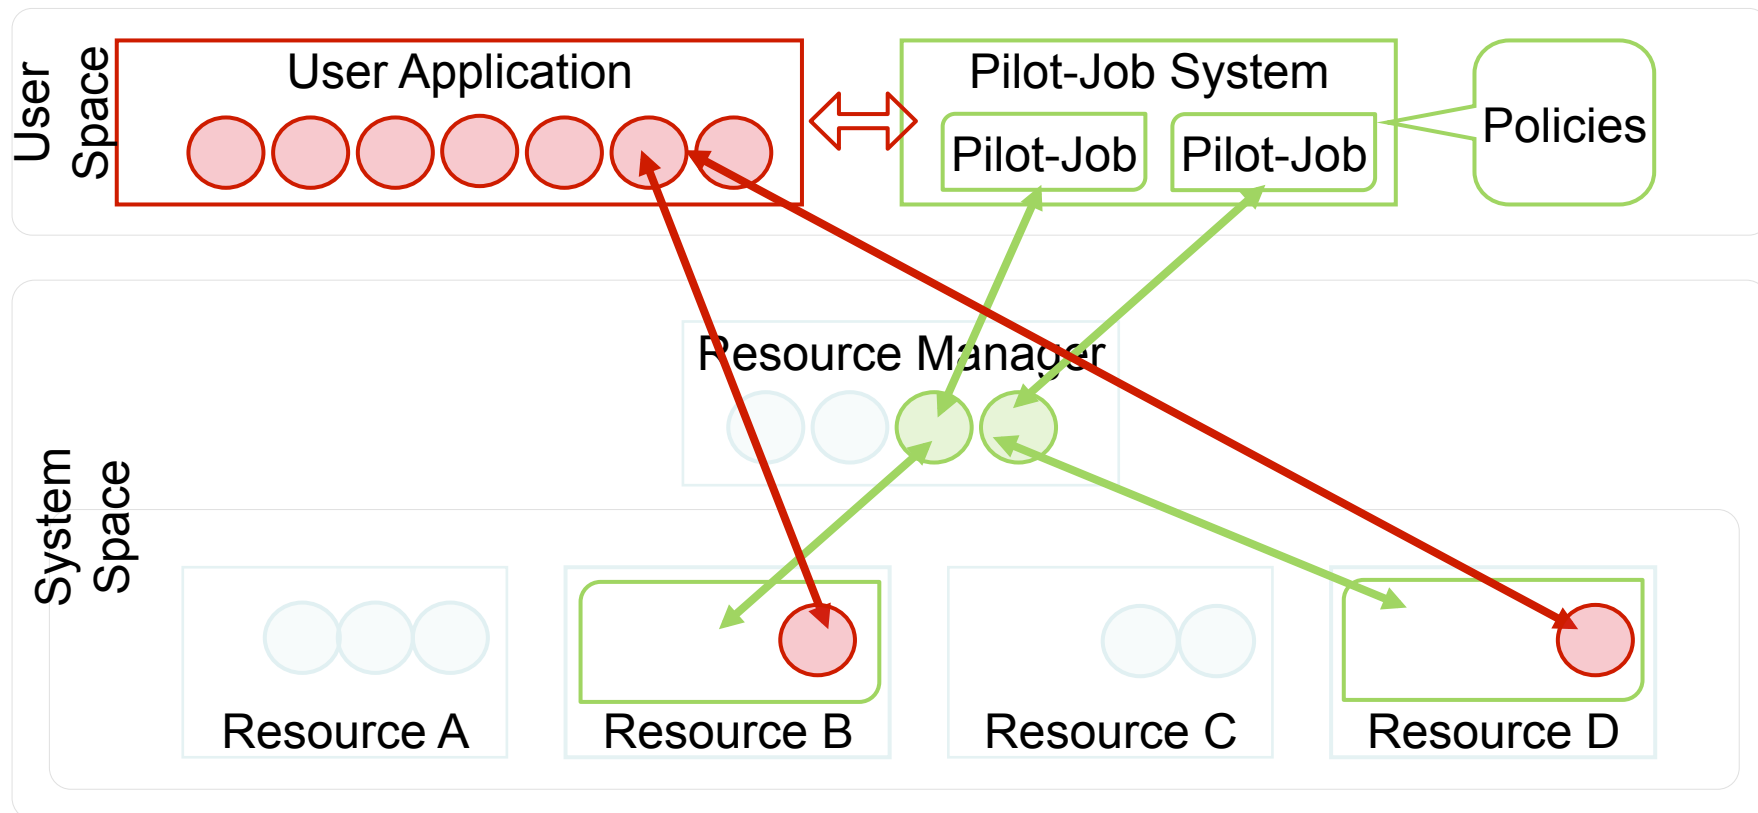
## PART-III

- Theory and Practice of Pilot Abstractions:
  - P\*: A Model for Pilot Abstractions
  - BigJob: An implementation of P\* Model
- Building higher-level capabilities on Pilot Abstractions
  - Pilot-MapReduce, Replica-Exchange

## PART-I: Introducing Pilot-Abstractions

# Introduction to Pilot-Abstraction

- **Working definition:** a system that generalizes a placeholder job to provide multi-level scheduling to allow application-level control over the system scheduler via a scheduling overlay



## Introduction to Pilot-Abstraction (2)

- **Working definitions:**
  - A system that generalizes a placeholder job to provide multi-level scheduling to allow application-level control over the system scheduler via a scheduling overlay
  - “.. defined as an **abstraction** that generalizes the reoccurring concept of **utilizing a placeholder job** as a container for a **set of compute tasks**; an instance of that placeholder job is referred to as *Pilot-Job* or *pilot*.”
- **Advantages** of Pilot-Abstractions:
  - The Perfect Pilot: Decouples workload from resource management
  - Flexible Resource Management
    - Enables the fine-grained (ie “slicing and dicing”) of resources
    - Tighter temporal control and other advantages of application-level Scheduling (avoid limitations of system-level only scheduling)
  - Move control, extensibility and flexibility “upwards”
    - Build higher-level capabilities without explicit resource management

## Pilot-Jobs Systems (PJS): Five Myths

- PJS do not need well defined architecture, model and semantics, or, PJs are such a simple concept, they don't need more “attention”
  - Not to confuse “simple to use” with simple to design”
- PJS are only about meta-scheduling (reducing queuing delays) on HTC, or, PJS unfairly game HPC queuing
  - There are interesting usage modes beyond “cycle stealing”
- PJS have to be tied to specific DCI; DCI are tied to specific PJ
  - Extensibility and interoperability have been difficult to establish
- ***PJS are passive (system) tools, as opposed to user-space, active and extensible components of a CI***
  - PJs can be user-controlled “programmable” elements
- PJS do not help with next-generation “data-intensive” applications
  - PJ for NGS O(10-100) GB per task on existing DCI



# Landscape of Pilot-Job Systems

- There are many PJS offerings, often semantically distinct
  - PanDA, DIANE, DIRAC, Condor Glide-In, SWIFT, ToPoS Falkon, BigJob...
    - *Why do you think there has been a proliferation of PJs?*
- Difference in the execution models of the PJ
  - We know “what” pilot-jobs do, but the “how” remains less clear
    - How to map tasks to pilot-jobs? How to choose/map optimal resource?
    - How to “slice and dice” resources?
- Conceptual & practical barriers to extensibility (& interoperability)
  - The landscape of PJS reflects, in addition to PJS specifics, the broader eco-system of distributed middleware & infrastructure
  - Software Engineering issues, interfaces, standardization
- Data remains a dependent variable, not a primary variable
  - Introduce the concept of Pilot-data

# Pilot-Job Paradigm

Based upon analysis of several Pilot-Job implementations

**Architecture:** Three distinct logical elements:

- **Workload Manager:** Responsible for making available the tasks to the executor alongside the needed data and retrieving results
- **Task Executor:** Responsible for executing the tasks while managing their data.
- **Communication and Coordination (C-C):** Patterns allow for and regulate the interaction between (and within) these two components.

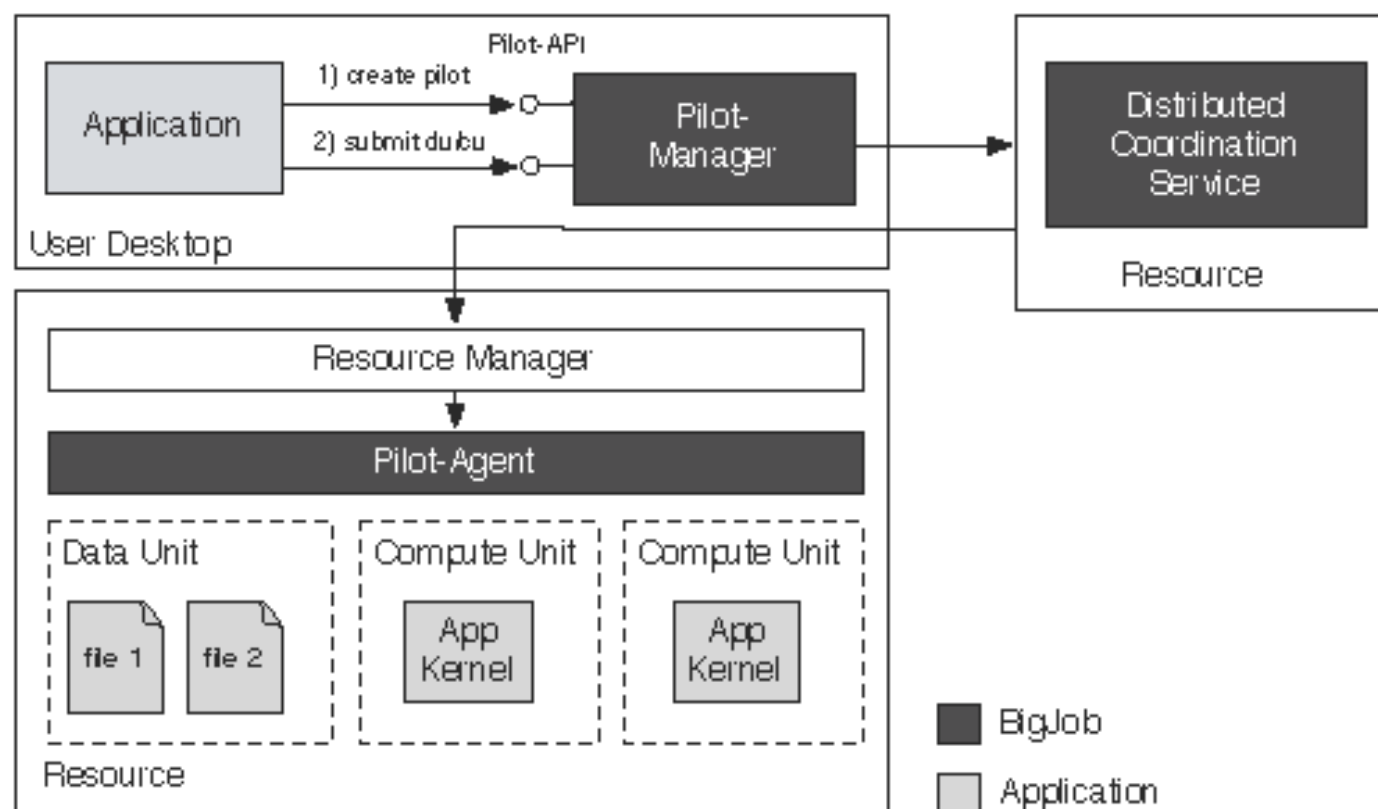
**Execution Patterns:** Based on multi-level scheduling and late-binding

- **Multi-level scheduling.** Tasks of a workload are scheduled on one or more pilots and the pilots are then scheduled on a given resource

**Capability/Functionality:** A system that generalizes a placeholder job to provide multi-level scheduling to allow application-level control over the system scheduler via a scheduling overlay

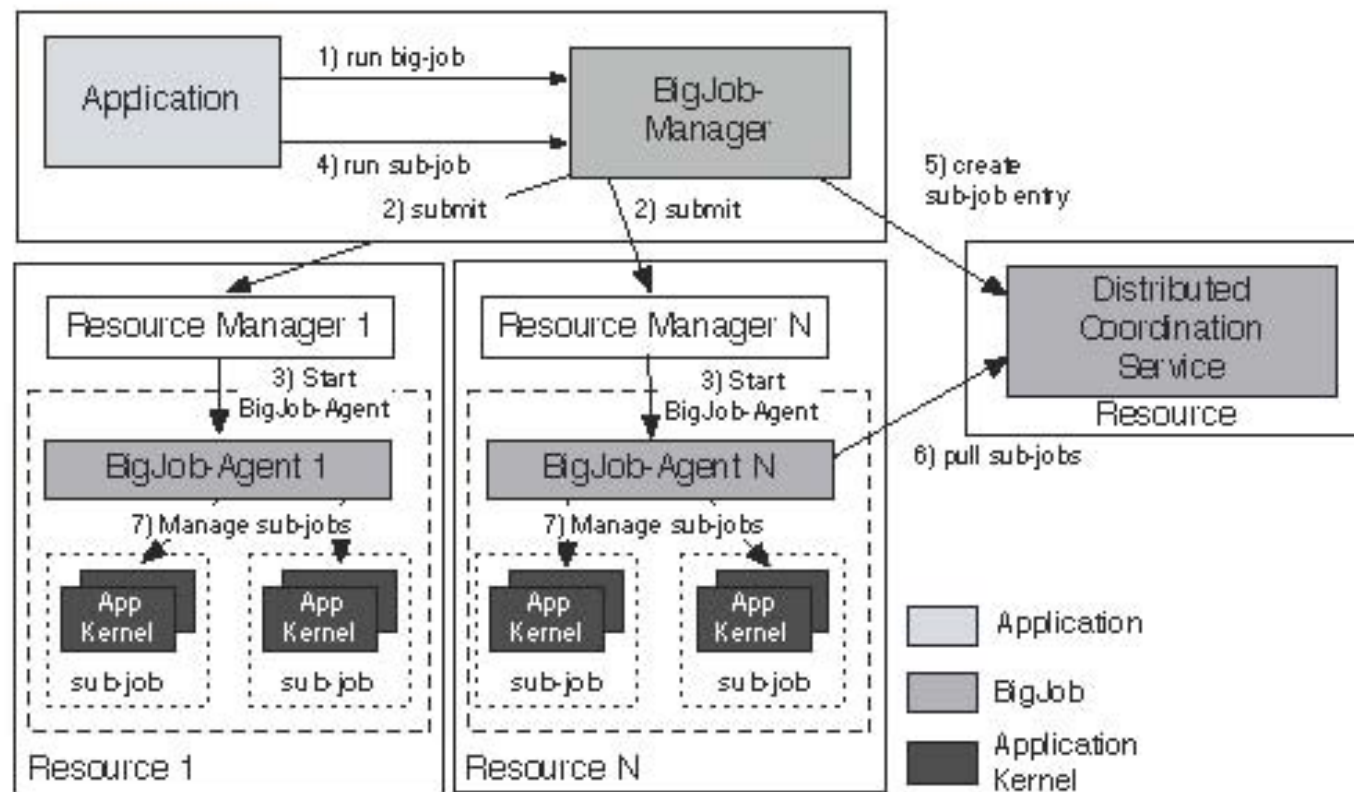
# BigJob: A simple, extensible and interoperable Pilot-Job System

# BigJob: Architecture

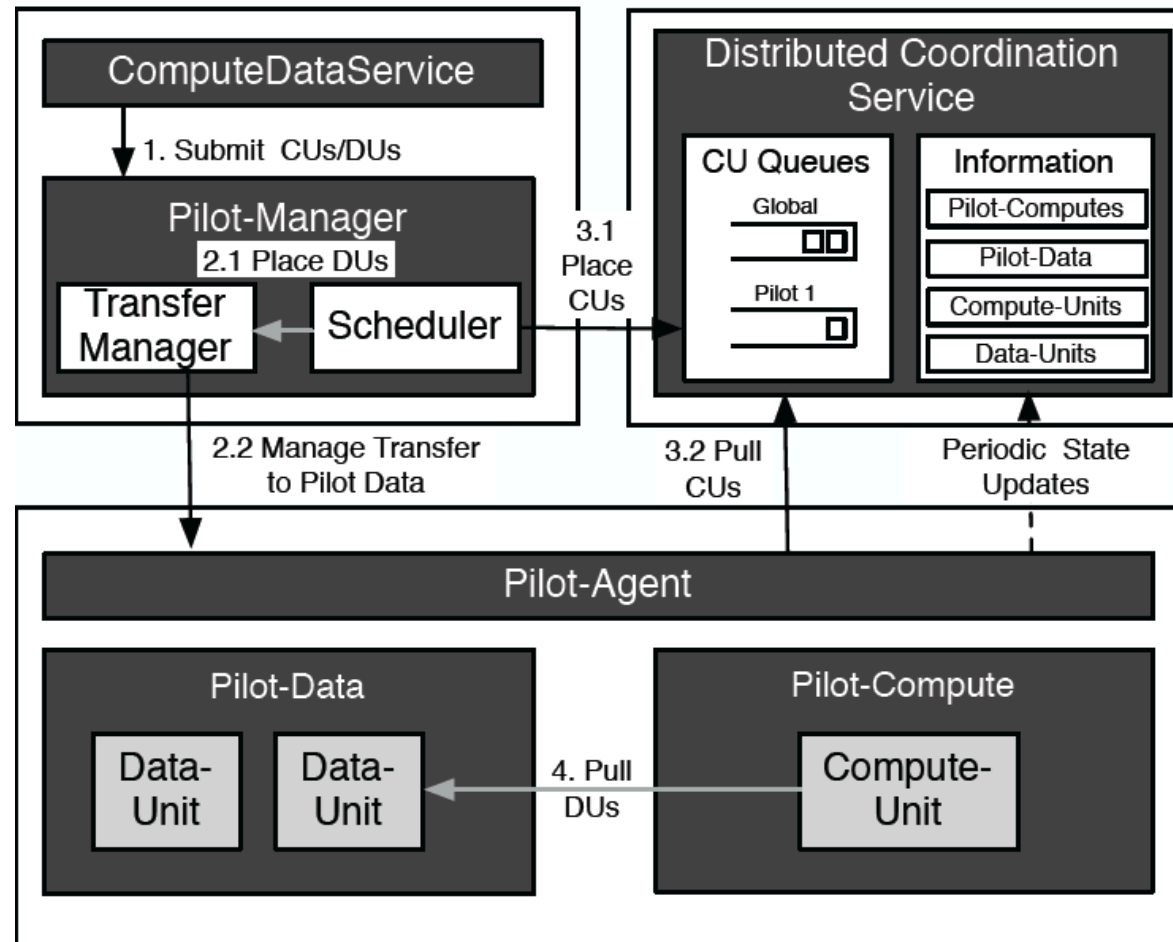


# BigJob: Distributed View

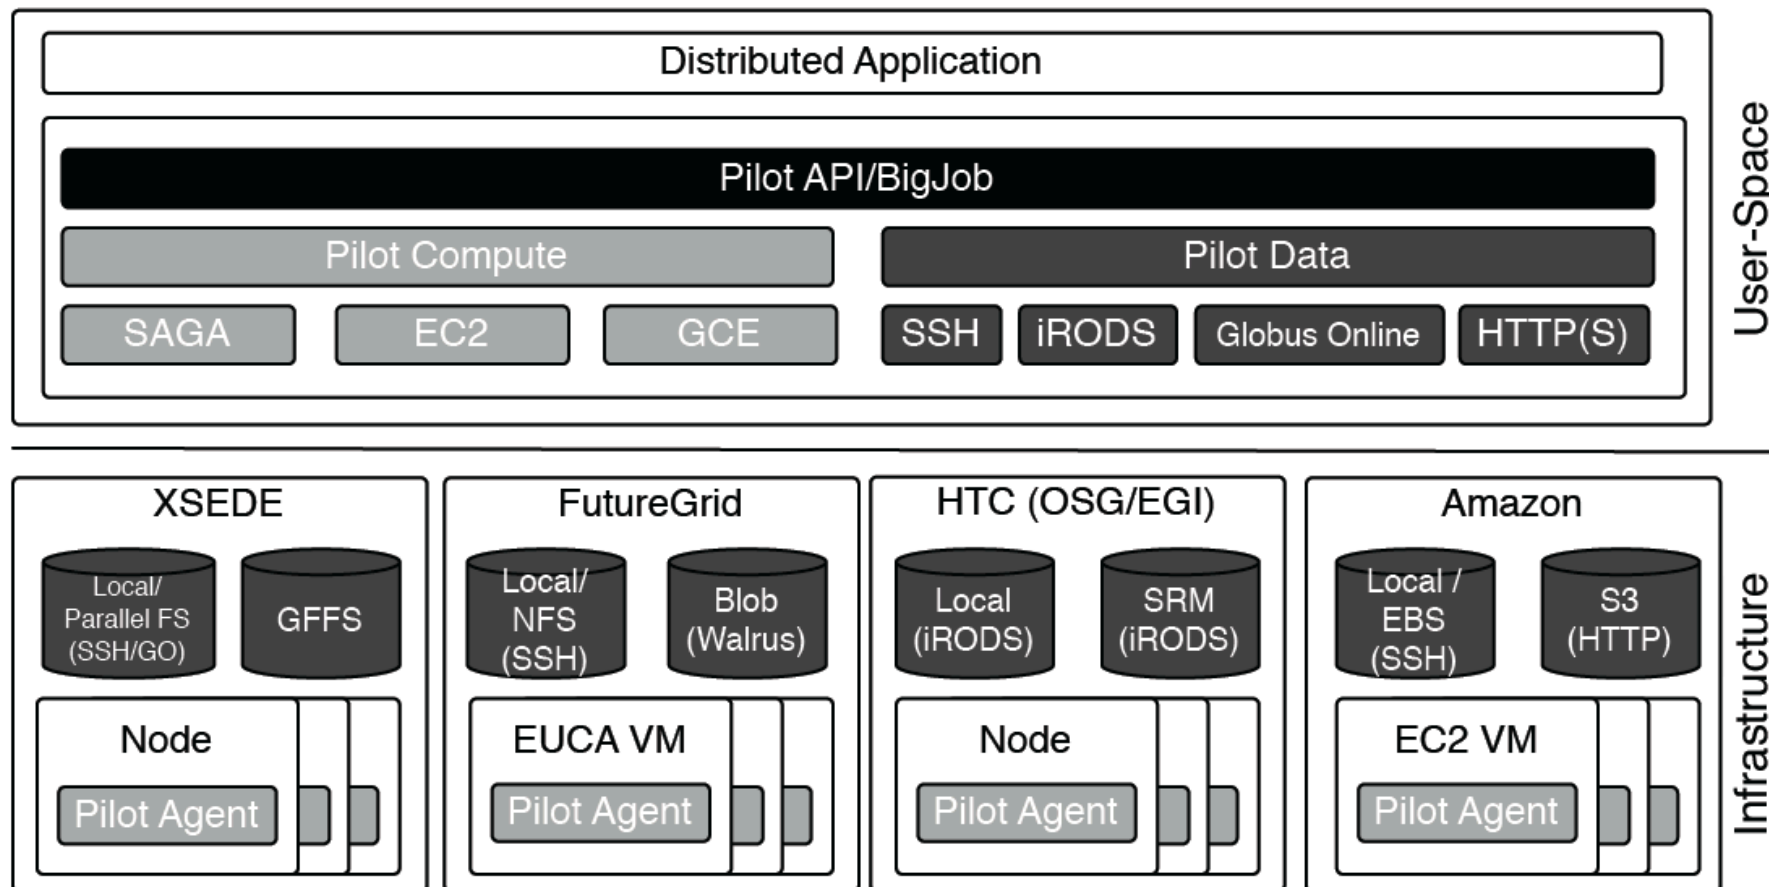
(Strongly encourage you to try after the tutorial)



# BigJob Workload Management

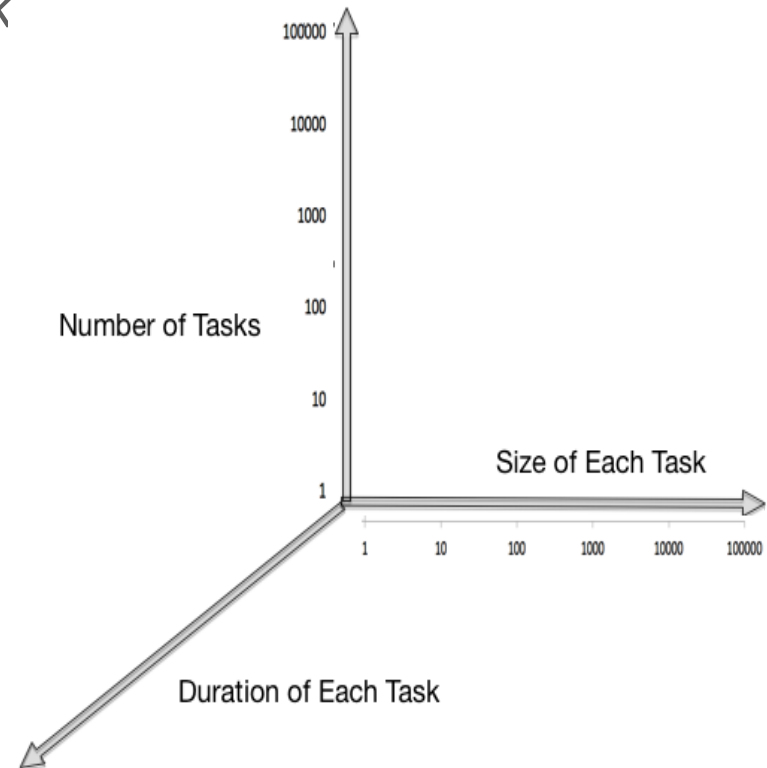


# BigJob: Resource Interoperability



## Design Objective: Scaling Along Many Dimensions

- Scaling Dimension #1: Size of each task
  - *Scale-up*
  - 1 core per task, 128/512.....
- Scaling Dimension #2: Total Number of Concurrent tasks
  - *Scale-Out*
  - Enhanced (MD) sampling  $O(1000)$ , statistical errors  $O(10^6)$  tasks,
- Scaling Dimension #3: Total Number of Tasks/per unit time
  - Real time processing  $O(1000)$
- Scaling Dimension #4: Number of Resources Used
  - *Scale-Across*
  - Execute on Grids, Clouds, Clusters







## “Coarse-Grained” BigJob Performance

- Number of zero-payload tasks that BJ can dispatch per second:
  - Distributed:  $O(10)$
  - Locally:  $> O(10)$
- Number of Pilots (Pilot-Agents) that can be marshaled
  - Locally/Distributed:  $O(100)$
- Typical number of tasks per Pilot-Agent:
  - Locally/distributed:  $O(1000)$
- Number of tasks concurrently managed = Number of Pilot-Agents x tasks per each agent :
  - $O(100) \times O(1000)$
- (Obviously) The above depends upon data per task:
  - BigJob has been used over  $O(1)$ -- $O(10^9)$  bytes/task, for tasks of duration  $O(1)$  second to  $O(10^5)$  seconds

# BigJob: (Partial) Usage on XSEDE Machines



> 10M SUs/year (and increasing) on XSEDE machines

# SAGA: Interoperability Layer for BigJob

<http://saga-project.org>

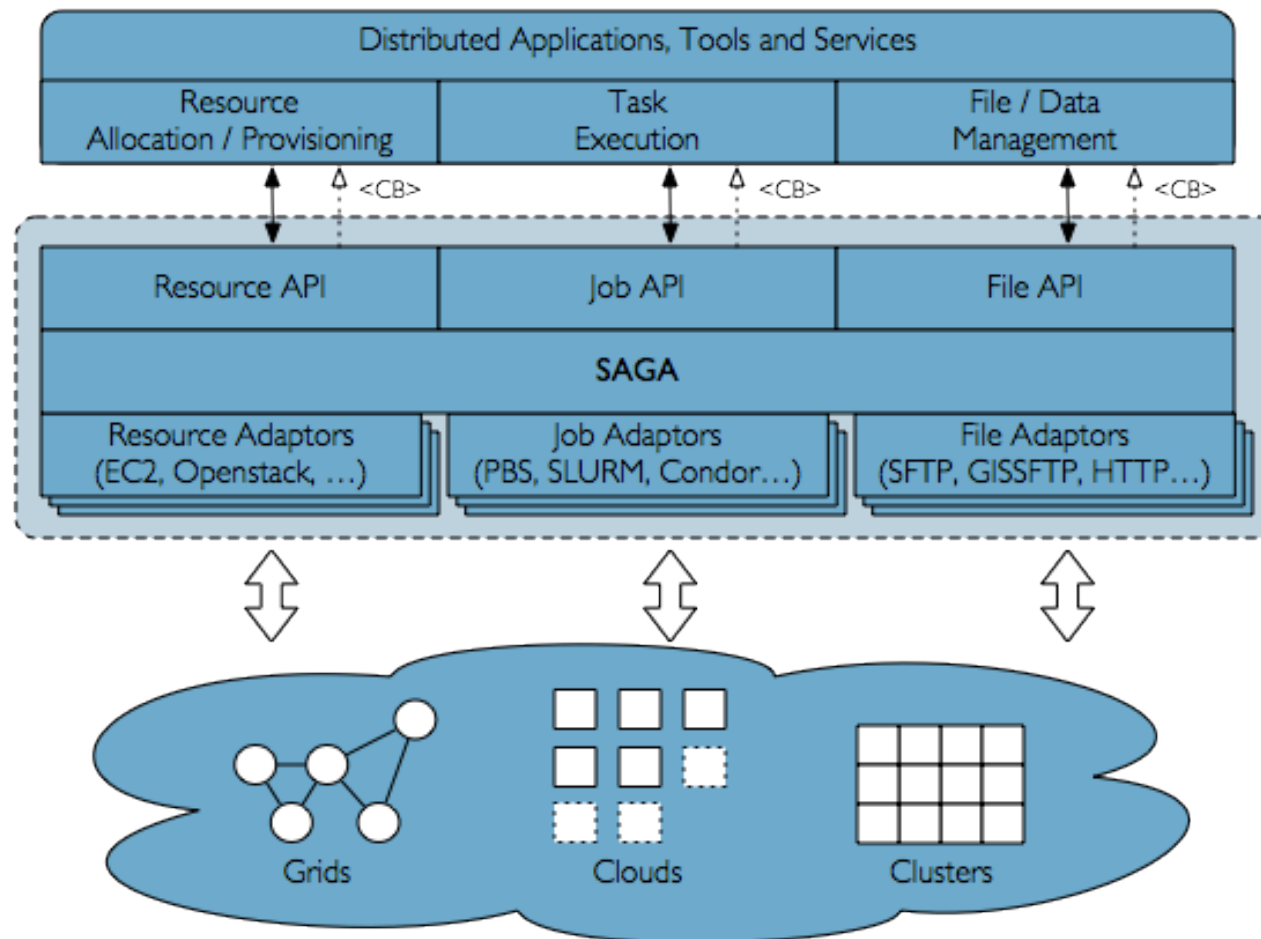
<http://saga-project.github.io/saga-python>

A Light-Weight Python Access Layer for Distributed Computing  
Infrastructure


# Overview

- SAGA: Simple API for Distributed (“Grid”) Applications
  - Application level standardized (Open Grid Forum GFD.90) API
  - Application is a broad term: “one person’s application is another person’s tool (building block)”.
- SAGA-Python:
  - Native Python implementation of Open Grid Forum GFD.90
  - Allows access to different middleware / services through a unified interface
  - Provides access via different backend plug-ins (“adaptors”)
  - SAGA-Python provides both a common API, but also unified semantics across heterogeneous middleware:
    - Transparent Remote operations (SSH / GSISsh tunneling)
    - Asynchronous operations
    - Callbacks
    - Error Handling

# Schematic

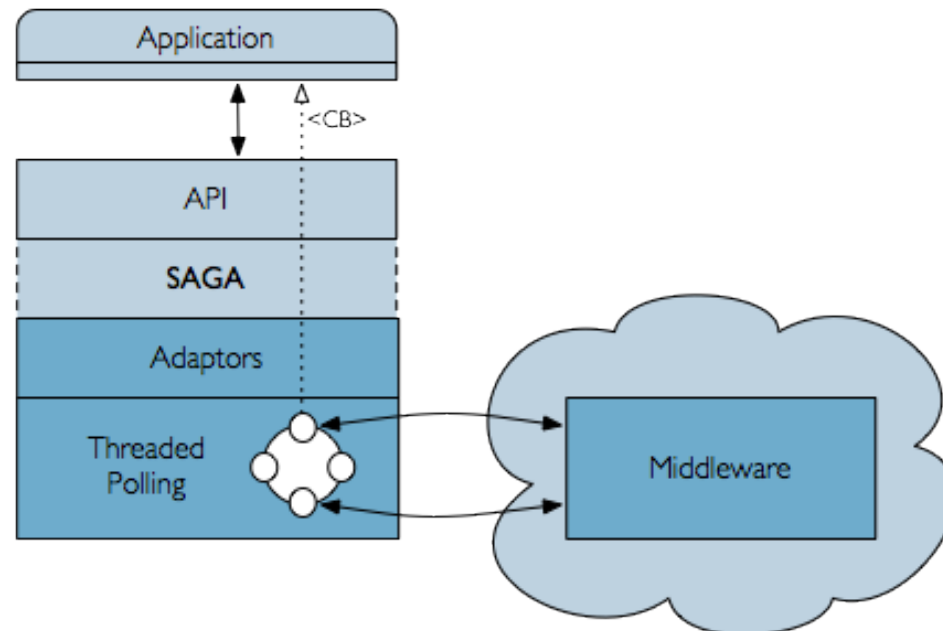


# Transparent Remote Operations



- Fast and optimized (GSI-)SSH command transport wrapper can be used by adaptors to access otherwise 'local-only' services, like many queuing systems

# Transparent CALLBACKS



- Callbacks and asynchronous operations are important concepts in distributed applications but are not supported by all middleware systems
- SAGA-Python moves application-level polling into the adaptor layer and exposes a clean callback interface through the API

# Available adaptors

<http://saga-project.github.io/saga-python/doc/adaptorssaga.adaptor.index.html>

- Job Submission Systems
  - SSH, GSISSH, Condor, Condor-G, PBS(-Pro), TORQUE, SGE, SLURM.
  - Under development: LSF
- File / Data Management
  - SFTP, GSIFTP, HTTP, HTTPS.
  - Under development: iRODS (Globus Online)
- Resource Management / Clouds
  - Amazon EC2, Openstack ('libcloud'-based)

[http://saga-project.github.io/saga-python/doc/adaptors/saga.adaptor.ec2\\_resource.html](http://saga-project.github.io/saga-python/doc/adaptors/saga.adaptor.ec2_resource.html)





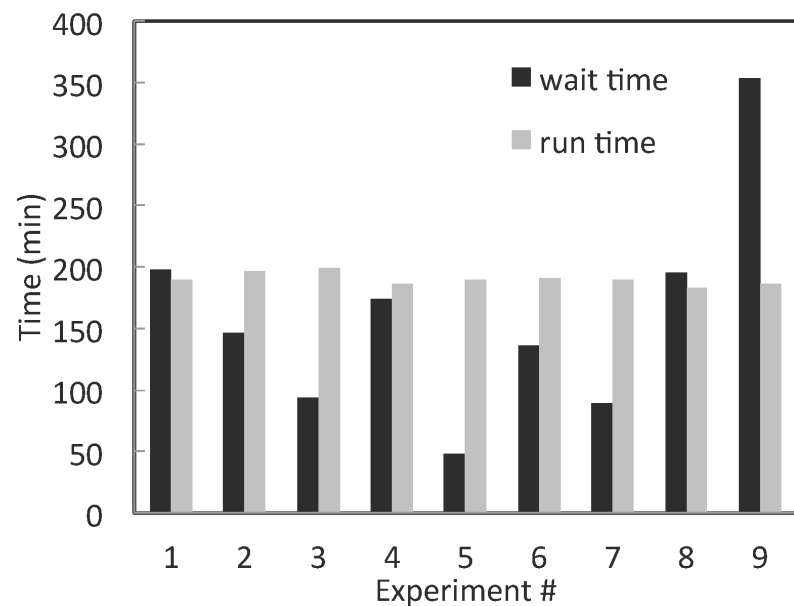
## SAGA: Interoperability layer upon which other tools and applications are built

- HOW SAGA is Used?
  - Uniform Access-layer to DCI, e.g, EGI, XSEDE, etc
  - Application “Scripting Layer” to DCI
  - Build tools, middleware services and capabilities e.g. Gateways,
- WHAT is SAGA Used for?
  - Production-grade science and engineering and research tool to design, implement and reason about distributed programming models, systems and applications
- Where is SAGA Used?
  - Pilot-job and Workflow systems, science gateways and web portals
  - Domain-specific (distributed) applications, libraries and frameworks

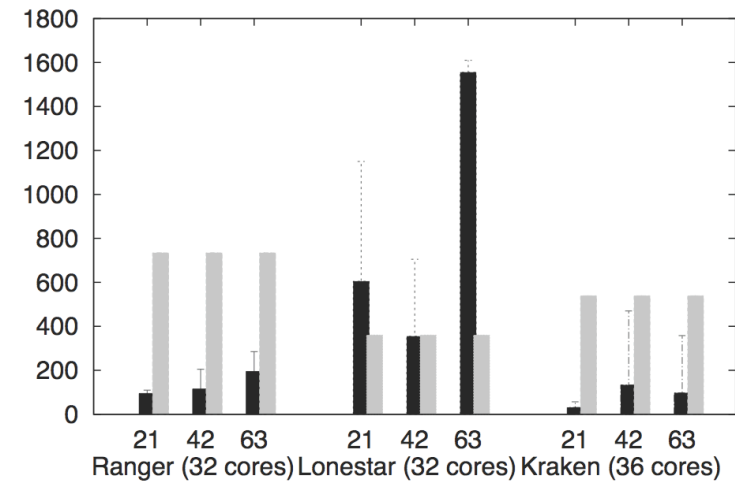
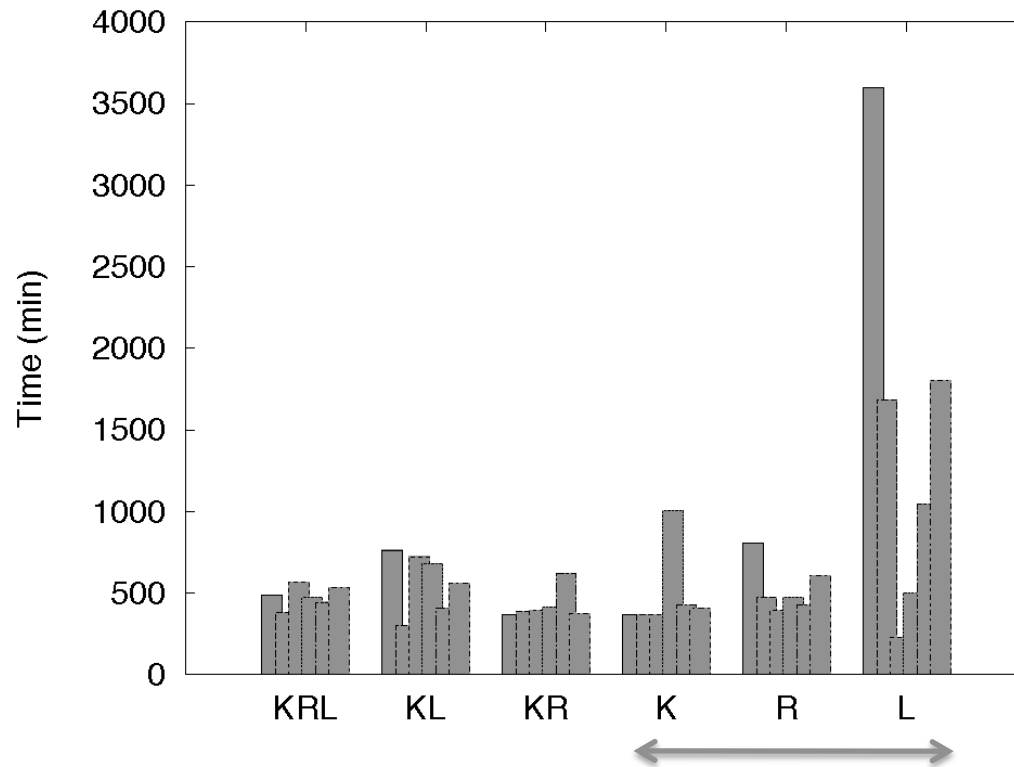
## Case Studies: Varying the “coupling” between many (MD) Simulations

# HT-HPC on Kraken

126 ensembles, each of 192 cores = 24192 cores

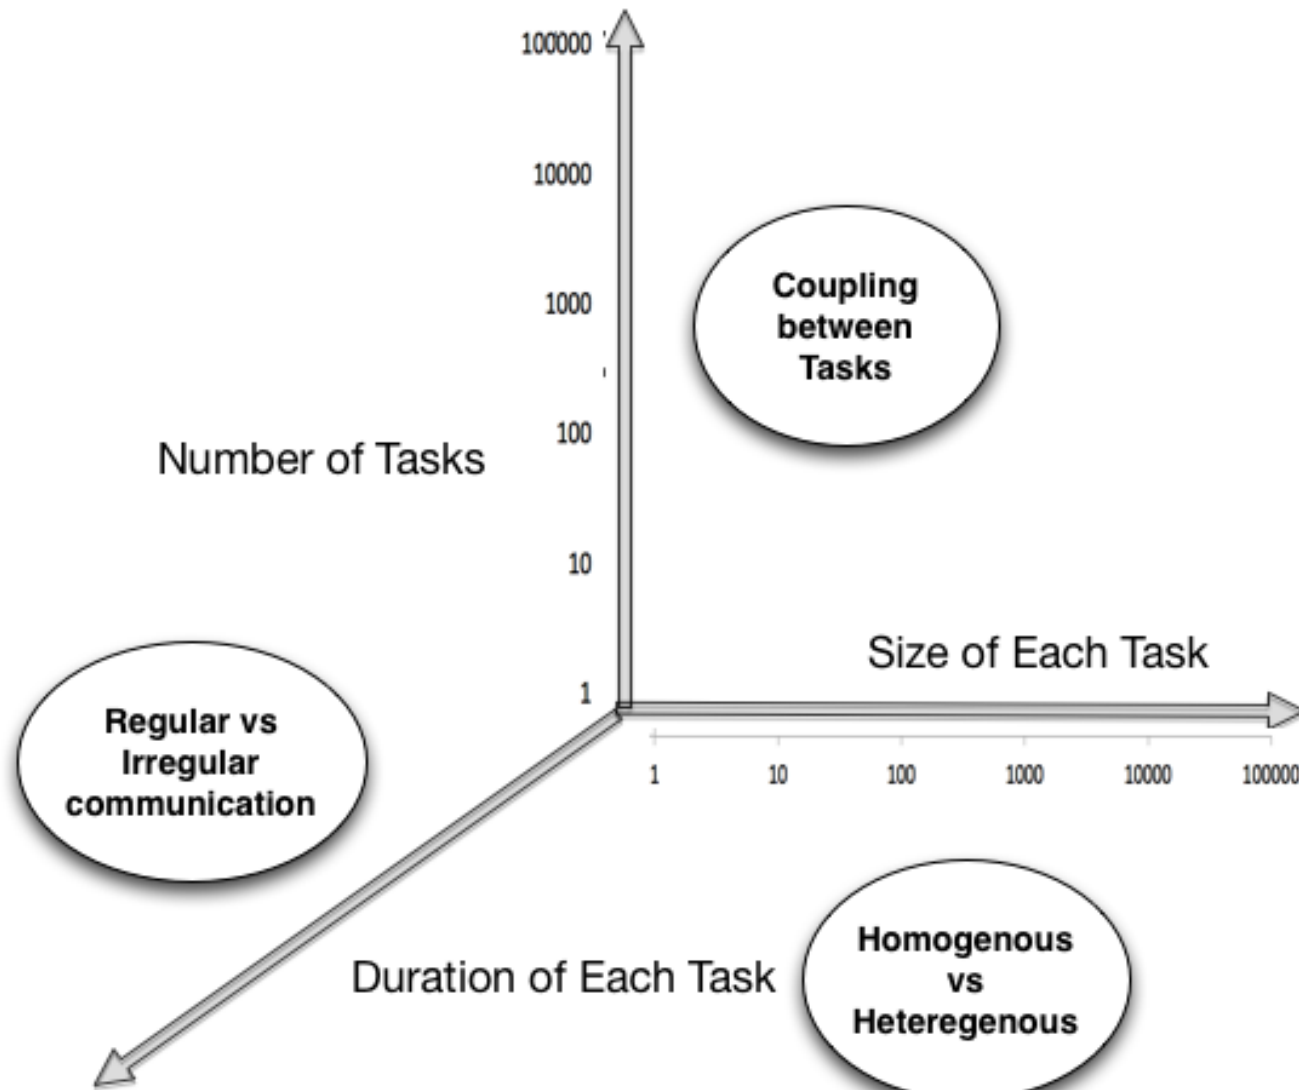


# Scale-Out/Across



X-axis: number of tasks (size)

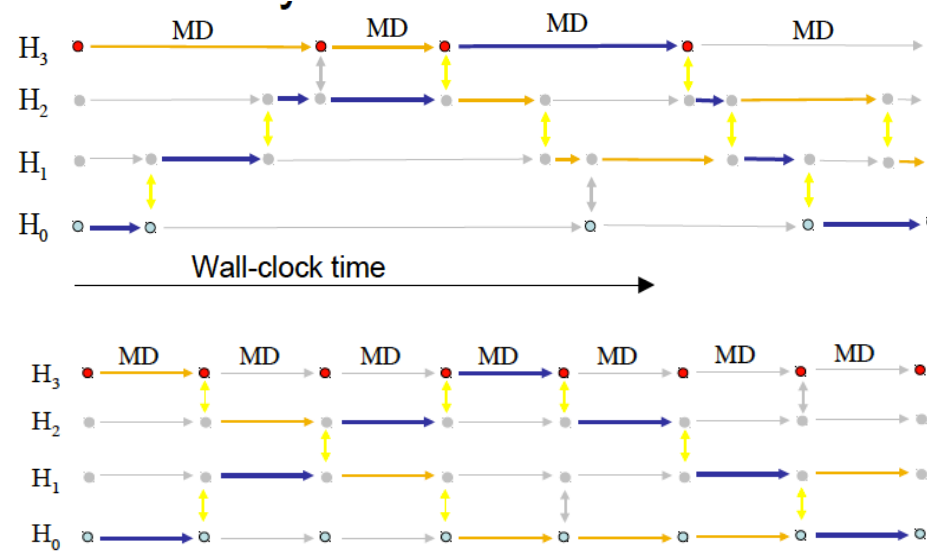
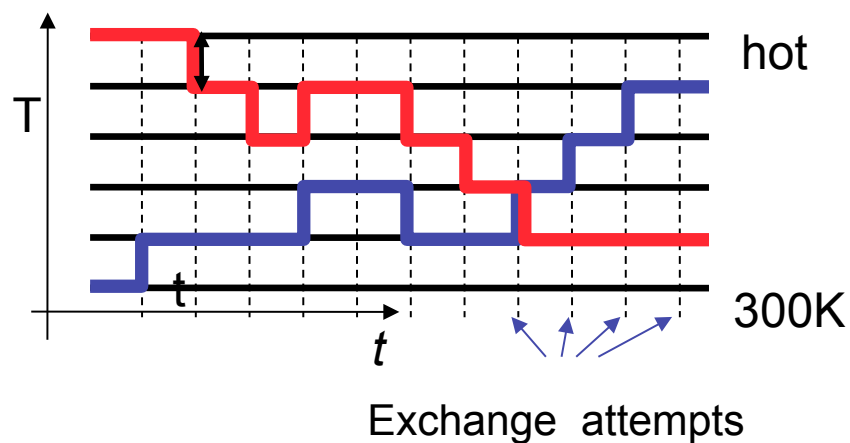


# Scaling Along Many Dimensions



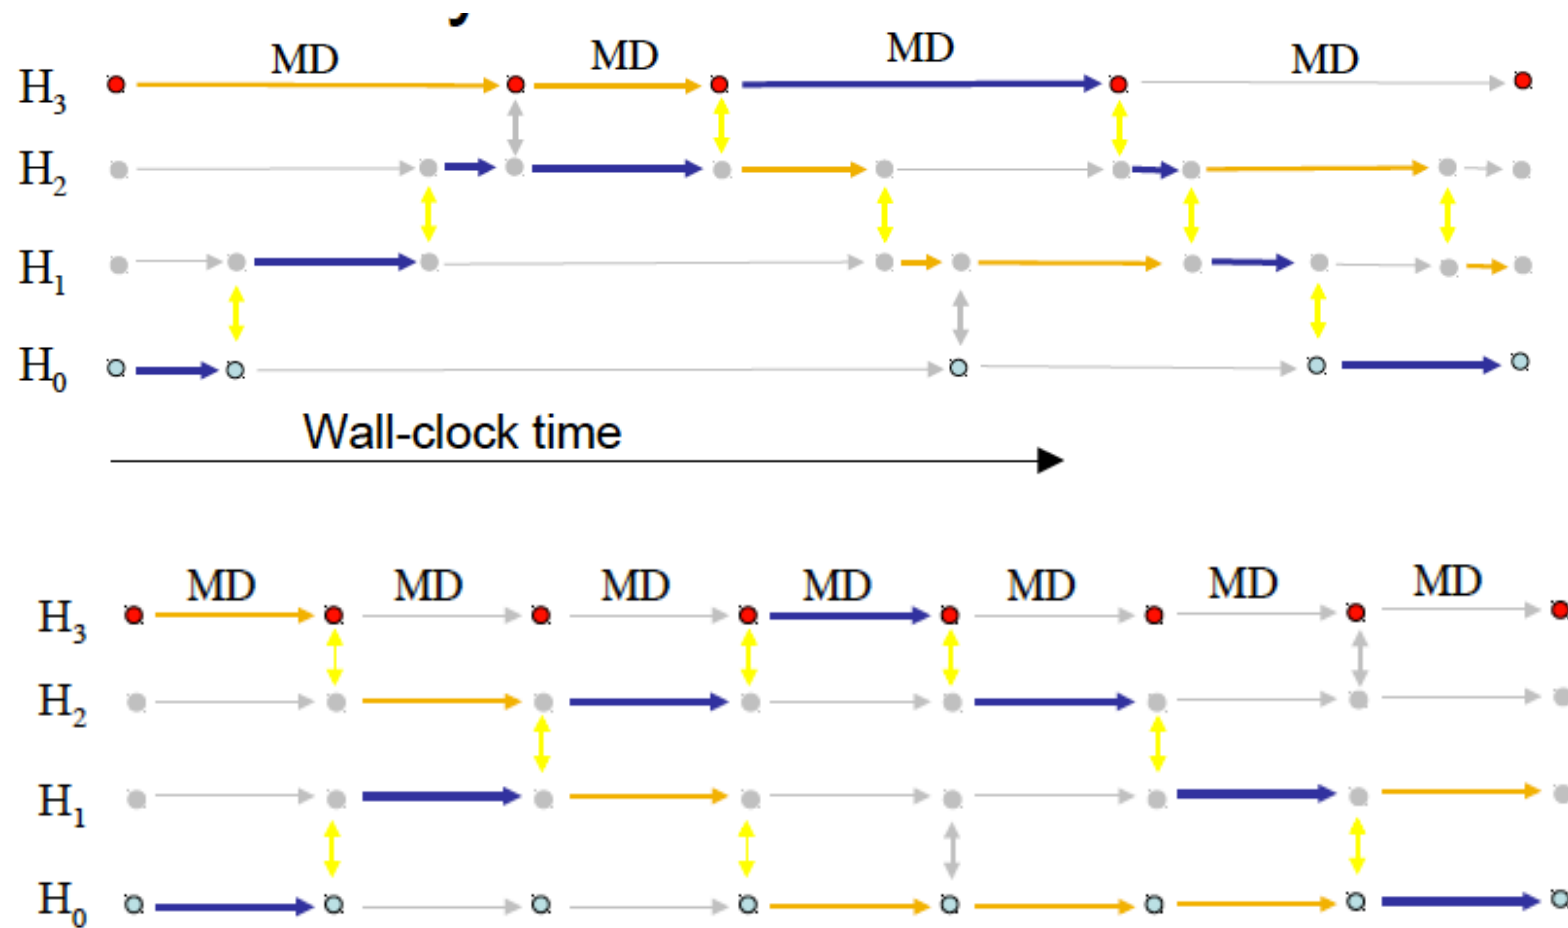
# Varying the Coupling between Simulations

- Large Ensemble of Uncoupled (Molecular) Simulations:
  - Enhanced Sampling of bio-molecular systems
- Multiple-Chaining of Simulations
  - Long time-scale simulations
- Ensemble of Loosely Coupled Simulations
  - Replica-Exchange Class of Algorithms
  - Understand protein-ligand recognition via (multidimensional) replica-exchange
- Requirements:
  - Launch and monitor/manage (order  $10^2$ - $10^4$ ) ensemble members
    - Where each ensemble member could be 128-1024 cores
  - Varying degrees of coupling between ensemble members
  - Varying job duration: hours to days to weeks
  - Ad hoc pair wise (a)synchronous data exchange

# Replica Exchange Molecular Simulations



# Irregular Synchronization





## PART II: HANDS ON TUTORIAL

# Getting Started

- These slides:

<https://github.com/saga-project/tutorials/blob/master/lecture/xsede13.pdf>

- URL for Tutorial:

<https://github.com/saga-project/tutorials/wiki/XSEDE13>

## PART III: PRINCIPLES AND PRACTISE OF PILOT ABSTRACTIONS

# PART III: PRINCIPLES AND PRACTISE OF PILOT ABSTRACTIONS

OR

*What makes BigJob Unique?*

# $P^*$ : A Model of Pilot-Abstractions

# P\*: A **Conceptual** Model for Pilot Abstractions

- A minimal but complete model
  - **Minimal:** Towards a common understanding of pilot-jobs
    - Provides vocabulary and model for analysis and insight across different PJ systems
    - Does not try to provide closed form answers to all issues
  - **Complete:** Can be used to design a priori an effective Pilot-Job
- Establish basic terminology, functionality and inter-relationship
  - What is the pilot? What is the agent? Push or pull?
- Provide a framework for comparison across Pilot implementations
  - Many existing pilot implementations can be understood using P\*
  - Not all pilot implementations must adhere strictly to P\*
- Unified view of Pilot-Abstractions
  - Provides symmetrical treatment for compute and data (and eventually network)

# P\* Model: Elements, Characteristics and API

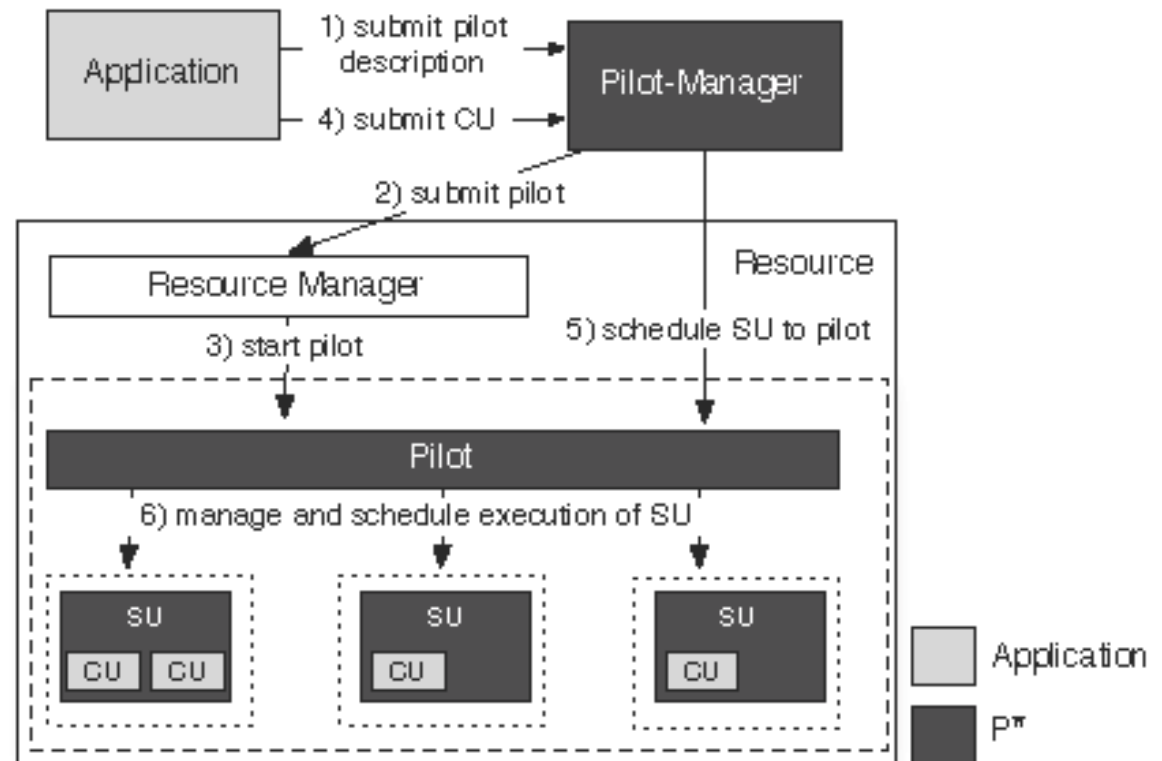
- **Elements:**

- Pilot-Compute (PC)
- Pilot-Data (PD)
- Compute Unit (CU)
- Data Unit (DU)
- Scheduling Unit (SU)
- Pilot-Manager (PM)

- **Characteristics:**

- Coordination
- Communication
- Scheduling

- **Pilot-API**



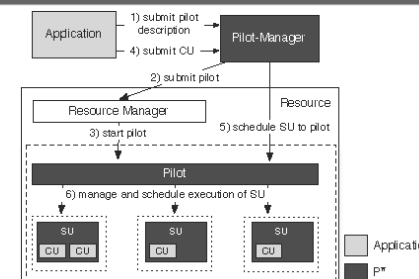
# P\* Elements

## Pilot-Compute

- The placeholder entity that gets submitted to a resource
- Also, associated with the role of an agent:
  - collects information
  - manages the resources allocated
  - exchanges data
- Executes application code

## Pilot-Data

- The placeholder entity that represents a storage resource (reservation)
- Can have the role of an agent:
  - collects information
  - manages the resources allocated
- Physically stores the data





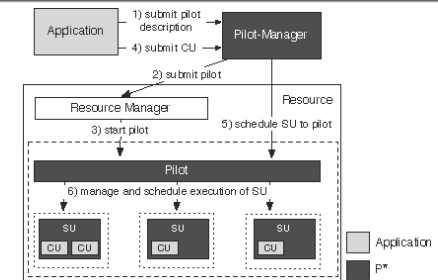
# P\* Elements

## Compute Unit (CU)

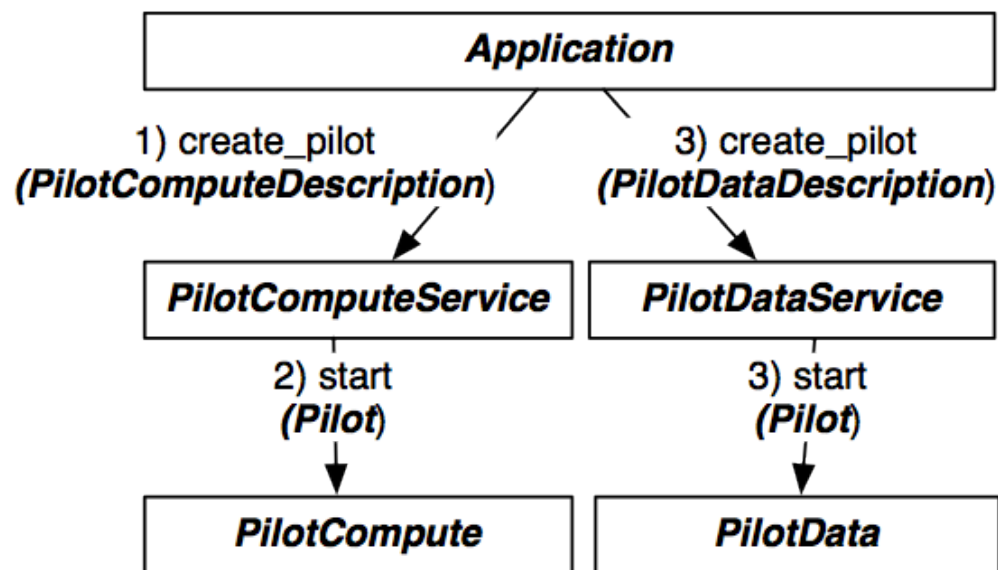
- Is defined by the application
- Encapsulates a self-contained piece of logical work that is submitted to the PJ system. E.g.:
  - task, job, rpc, web service call, etc.

## Data Unit (DU)

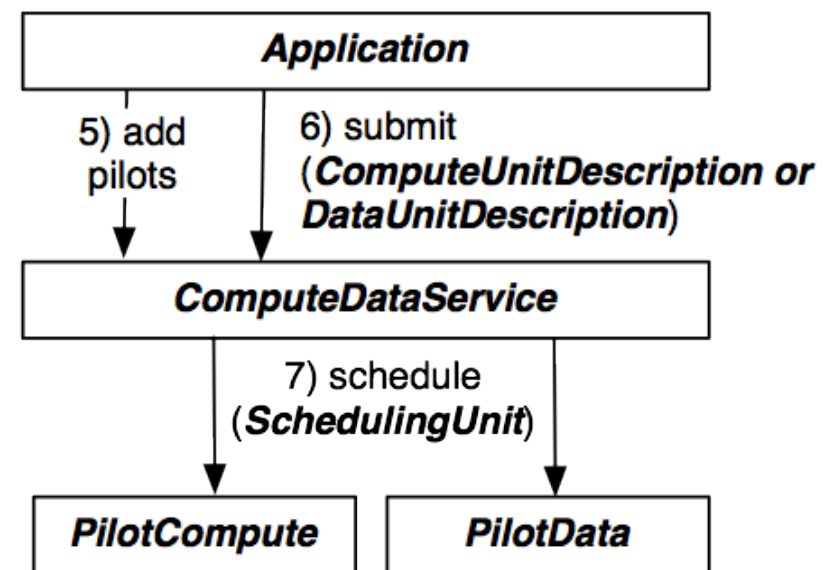
- Is defined by the application
- Encapsulates a self-contained piece of logical data that is submitted to the PJ system. E.g.:
  - file, chunk, database, etc.



# Pilot-API: Unified API to Pilot-Compute and Pilot-Data



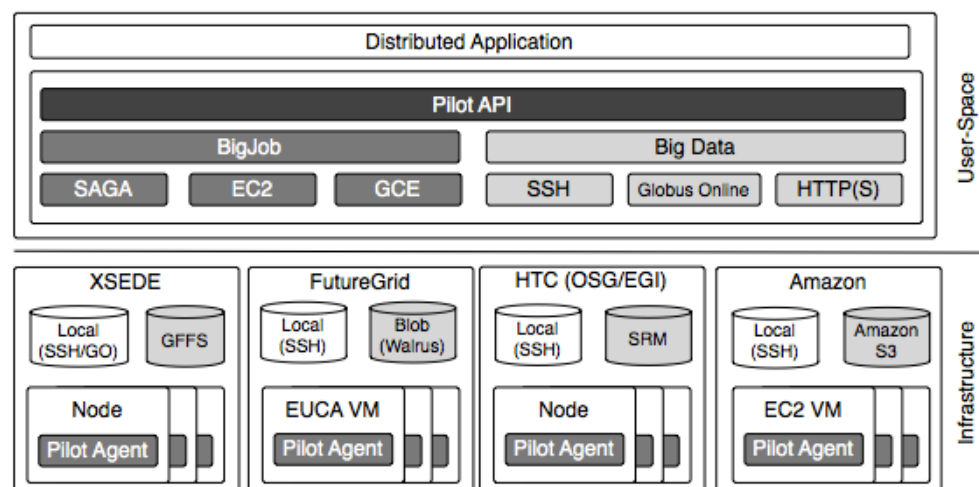
Managing Pilots



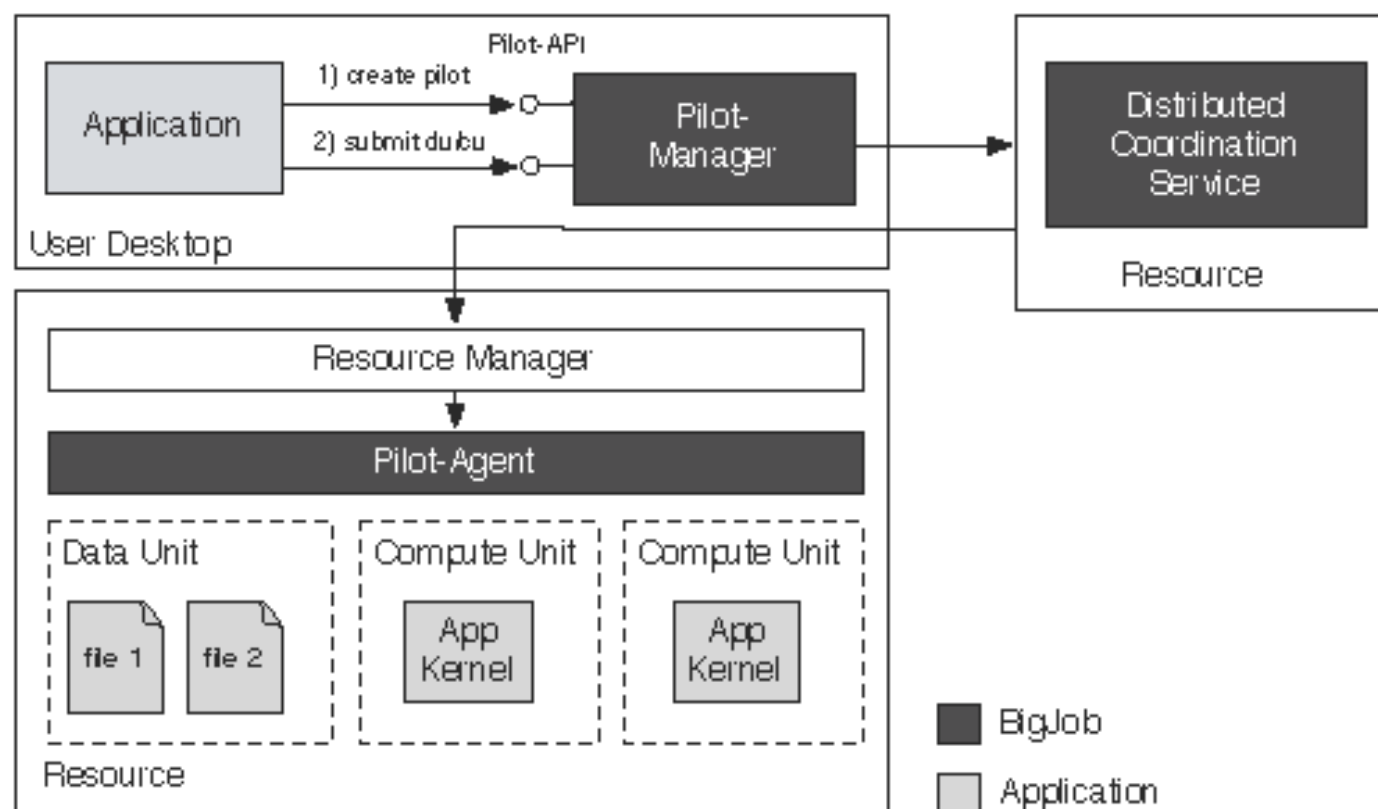
Managing Application Workload

# P\*: Mapping and Interoperability

P* Element		BigJob	DIANE	Condor-G/ Glide-in
Pilot-Manager		BigJob Manager	RunMaster	condor_master condor_collector condor_negotiator condor_schedd
Pilot		BigJob Agent	Worker Agent	condor_master condor_startd
Compute Unit (CU)		Task	Task	Job
Scheduling (SU)	Unit	Sub-Job	Task	Job



# BigJob: Architecture



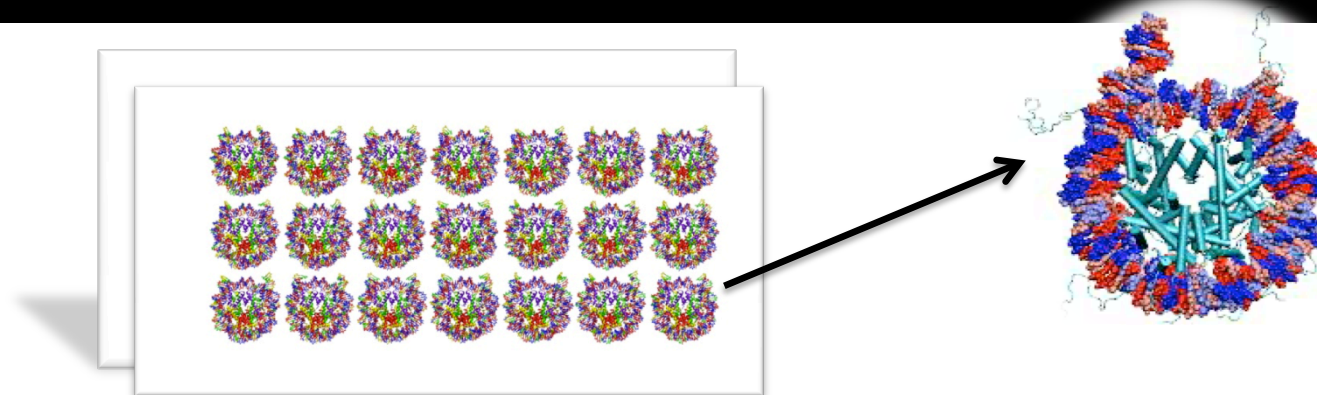
# BigJob @ XSEDE'13

Talk by Jack Smith @ 3:45pm Tuesday

Talk by Melissa Romanus: 11am Wednesday

Poster Session: Jacob Swadling

# Scalable Online Comparative Genomics of Mononucleosomes




Collaboration with Tom Bishop, Jack Smith  
and Work supported by XSEDE Fellows Grant (Jack Smith)  
See Talk by Jack Smith @ 3:45pm Tuesday

# Mapping Complex Biomolecular Reactions using Large-Scale Replica-Exchange Simulations on National Production Cyberinfrastructure

Collaboration with Darrin York, Ron Levy  
Talk by Melissa Romanus: 11am Wednesday


# Replica-Exchange Performance in AMBER



Umbrella sampling of the backbone conformational space of alanine dipeptide. The exchange parameters are all permutations of harmonic biasing potentials on each torsion.



# QM/MM Replica-Exchange



- QM/MM is not generally considered parallelizable
  - RE provides efficient sampling irrespective
- Repex FW: BigJob-based Repex
  - Framework to handle  $O(100)$ ---  $O(1000)$  concurrent simulations
  - Amongst the earliest QM/MM
  - Framework similar to classical QM/MM
- Performance increases linearly with core count and number of QM/MM simulations
  - No apparent coordination cost for additional simulations

## Pilot-Mapreduce: Library for MR using Pilot-API for Dynamic Resource Management

# BigJob Summary

- BigJob: Available Simple, Extensible and Interoperable PJS
- Simple: Adheres to a simple conceptual model of pilot-abstractions
  - Well defined sub-systems, execution model, semantics
  - We now know both “what” the PJS does, but also the “how”!
    - How to map tasks to pilots?
    - How to “slice and dice” resources?
- Extensible:
  - User controlled
  - Programmable (via Pilot-API)
  - Different application types, execution modes, coupling schemes
- Interoperable: Across XSEDE resources

## (More general) Conclusions

- “Pilot Abstraction Works!”
  - Abstraction for separation of workload and resource management,
  - Dynamic resource utilization
- SAGA-BigJob: A Pilot-Job System that is more than a Pilot-Job System!
  - Libraries for Replica-Exchange, PMR, Binding Affinity Calculator builds upon Pilots for (dynamic) resource management

# Acknowledgements

## **Graduate Students:**

- Ashley Zebrowski
- Melissa Romanus
- Mark Santcroos
- Anton Trekalis

## **Undergraduate Students:**

- Vishal Shah

## **Research Scientists:**

- Andre Luckow
- Andre Merzky
- Matteo Turilli
- Ole Weidner

- Anjani Raghotham
- Dinesh Prasad

## **TACC Staff**

- Yaakoub El-Khamra

## **Users of BigJob**

## **National Science Foundation**

## References

- **“P\*: A Model of Pilot-Abstractions”**, 8<sup>th</sup> IEEE International Conference on e-Science 2012 (DOI: 10.1109/eScience.2012.6404423)
- **“Distributed Computing Practice for Large-Scale Science & Engineering Applications”** Computing and Concurrency: Practice and Experience, 2012 (DOI: 10.1002/cpe.2897)
- **“Pilot-Data: An Abstraction for Distributed Data”**, arXiv:1301.6228
- **“Pilot MapReduce”**, 3<sup>rd</sup> Workshop on MapReduce and its applications (2012), in conjunction with HPDC (Delft)

## References

- SAGA-Python:
  - <http://saga-project.github.io/saga-python/>
- BigJob: An implementation of P\*
  - <http://saga-project.github.io/BigJob/>
- RADICAL:
  - <http://radical.rutgers.edu/>
- Publications:
  - <http://radical.rutgers.edu/publications>
- Tutorials:
  - <https://github.com/saga-project/tutorials/wiki/XSEDE13>